

# Generating Efficient Interpreters

Marcel Taubert (mt652)  
20962335

2023  
School of Computing  
MSc Advanced Computer Science  
University of Kent

Supervisor: Dr. Michael Vollmer

Total word count: 7003

September 5, 2023

## UNIVERSITY OF KENT

### ACCESS TO A MASTER'S DEGREE OR POSTGRADUATE DIPLOMA DISSERTATION

In accordance with the Regulations, I hereby confirm that I shall permit general access to my dissertation at the discretion of the University Librarian. I agree that copies of my dissertation may be made for Libraries and research workers on the understanding that no publication in any form is made of the contents without my permission.

**Notes for Candidates:** by submitting your dissertation, you agree with the following:

- 1 Where the examiners consider the dissertation to be of distinction standard, one copy may be deposited in the University Library and/or uploaded into Moodle as an example of good dissertation for future students.
- 2 If a copy is sent to the Library, it becomes the property of the University Library. The copyright in its contents remains with the candidate. A duplicated sheet is pasted into the front of every thesis or dissertation deposited in the Library. The wording on the sheet is:

*"I undertake not to use any matter contained in this thesis for publication in any form without the prior knowledge of the author."*

Every reader of the dissertation must sign and date this sheet.

- 3 The University has the right to publish the title of the dissertation and the abstract and to authorise others to do so.

M. Taubert

05.09.2023

SIGNATURE

DATE

Marcel Taubert

FULL NAME

(Please print, underlining surname, in order to assist the cataloguing of theses deposited in the Library.)

---

---

### CERTIFICATE ON SUBMISSION OF DISSERTATION

I certify that:

- 1 I have read the University Degree Regulations under which this submission is made;
- 2 In so far as the dissertation involves any collaborative research, the extent of this collaboration has been clearly indicated; and that any material which has been previously presented and accepted for the award of an academic qualification at this University or elsewhere has also been clearly identified in the dissertation.

M. Taubert

05.09.2023

SIGNATURE

DATE

This form should be completed and included in the submission of your dissertation.

# Contents

<b>1</b>	<b>Introduction</b>	<b>5</b>
1.1	What is an interpreter? . . . . .	5
1.2	What is a virtual machine interpreter? . . . . .	5
1.3	Virtual machine . . . . .	6
1.4	Optimizations . . . . .	6
1.5	Automation . . . . .	6
1.6	Objectives . . . . .	6
<b>2</b>	<b>Description of the problem</b>	<b>7</b>
2.1	Efficiency . . . . .	7
2.2	Automation . . . . .	7
<b>3</b>	<b>Literature review</b>	<b>8</b>
3.1	Programming Language . . . . .	8
3.2	Interpreter approaches . . . . .	9
3.3	Virtual Machines . . . . .	9
3.4	Output Format . . . . .	12
<b>4</b>	<b>Approaches</b>	<b>12</b>
<b>5</b>	<b>Description of work</b>	<b>14</b>
5.1	Tools used . . . . .	14
5.2	Virtual Machine . . . . .	14
5.2.1	Implementation . . . . .	14
5.2.2	Stack . . . . .	15
5.2.3	Byte code . . . . .	15
5.3	Interpreter . . . . .	16
5.3.1	Scanner . . . . .	16
5.3.2	Parser . . . . .	17
5.3.3	Interpreter (testing) . . . . .	17
5.4	Optimizations . . . . .	18
5.4.1	Replacing superinstructions in the byte code . . . . .	21
5.5	Input language . . . . .	21
<b>6</b>	<b>Results &amp; Benchmarks</b>	<b>23</b>
6.1	Setup . . . . .	23
6.2	Test Programs . . . . .	24
6.3	Results . . . . .	27
<b>7</b>	<b>Summary</b>	<b>30</b>
<b>8</b>	<b>Future Work</b>	<b>31</b>

## ***Abstract***

Building a virtual machine interpreter following the traditional approach of using C or assembly as an implementation language can be frustrating, furthermore since we have modern high-level languages that provides us with safety features and an overall better development experience. This paper describes the process of building a virtual machine interpreter in Rust which highlight the convenience of modern high-level language features. It shows the impact of optimizations and efficient implementations in visual benchmarks and compares against industry-standard approaches. We discussed the advantages and disadvantages of different implementation approaches such as stack-based and register-based virtual machines, compiler optimizations like superinstructions and constant folding or interpreter based optimizations such as the computed goto optimization.

# 1 Introduction

In the modern world we live in today, everyone of us uses technology which is powered by code. Technology is in our phones and computers and now even in smart light bulbs and smart toasters. The code running on our devices is written by programmers around the world but many people, including a substantial number of programmers, lack an understanding of how the code is executed on their devices. In this paper we will explore one approach to executing code called interpreting. More precisely we will explore the implementation of a virtual machine interpreter in the programming language Rust.

## 1.1 What is an interpreter?

When we want to execute our program written in a high-level language such as Python, we first need to translate the human-readable source coding into a format the computer can understand. We differentiate on two approaches how this can be achieved.

The first technique describes the process of code being compiled (translated) into machine code. The end product is a stand-alone program that can be executed at any time on the architecture it was compiled for.

The second approach is called 'interpreting'. There are two primary methods for how an interpreter works. One approach is called a tree-walk interpreter. In this case, the interpreter just walks the generated abstract syntax tree and executes it. The second approach is a virtual machine interpreter which involves an extra step between the generation of the syntax tree and the execution. A compiler walks the tree and generates byte code. This bytecode will then be fed into the virtual machine which executes it. We will use a virtual machine interpreter in this paper. Interpreters are generally easier to implement and have some other advantages that make them more approachable than compilers like portability or a fast edit-compile-run cycle as stated by the authors of the paper 'vmgen' [6].

## 1.2 What is a virtual machine interpreter?

A famous technique to implement interpreters is to build a virtual machine interpreter. A VM interpreter is generally divided into two systems. A frontend and a backend. [6]

The frontend consists of a compiler that takes the written code and produces a sequence of byte code instructions. The backend is a virtual machine that gets the stream of bytecode instructions as input and executes them. [6] The byte code or intermediate representation (IR) used in a VM interpreter is usually designed to be very similar to a real machine. [6] Some real-world examples of virtual machine interpreters are for example the JVM (Java Virtual Machine) or the PVM (Python Virtual Machine).

### 1.3 Virtual machine

When we are talking about virtual machines we distinguish between stack-based VMs and register-based VMs. Each of the versions has its advantages and disadvantages. Stack-based virtual machines are generally easier to implement than the register-based approach. That shows for example the complexity of the intermediate representation used in the virtual machine. The stack-based bytecode tends towards being smaller and by that more efficient in comparison to the rather complex bytecode instructions of the register-based approach.

typical stack based instruction:

```
Push 1    -- push value to the stack
```

typical register based instruction:

```
LD R1, 42 -- load value '42' into register R1
```

This simplicity is the reason we have decided to rely on a stack-based virtual machine in this paper.

### 1.4 Optimizations

Optimizations are essential for delivering efficient and performant execution of programs. These optimizations can be implemented by the compiler directly or by other parts of a VM interpreter such as the JIT (Just in Time) compiler, which are more complex and will not be covered in this paper. Optimizations generally involve a trade-off in compilation time and memory against execution speed. This results in additional passes on the generated byte code for example that are needed to implement certain optimizations. During the development of this project, we have looked at several different optimizations. Some of the optimizations that we have implemented for the project are threaded code or the computed goto optimization inside of the interpreter and superinstructions and constant folding in the compiler.

### 1.5 Automation

Writing an interpreter for a programming language can be a tedious and challenging task on its own. Thinking about how to keep the code efficient and implement optimizations makes it even harder. One solution to this is automating the process of writing a virtual machine interpreter. This would prevent many error sources such as human errors and lead to enhanced code quality. By eliminating the manual writing work, the programmer can focus on higher-level aspects of the project. Due to quicker development cycles, the interpreter can be fine-tuned without needing to change the code and be more efficient.

### 1.6 Objectives

This paper will use the Rust programming language to build each part of a virtual machine interpreter with the eventual goal of automating the generation

of itself depending on a given configuration. We will explore techniques and approaches on how to build a VM interpreter in Rust and use benchmarking to visualize results, highlighting the effects of optimizations. Additionally, we will compare the results obtained using Rust's convenient high-level features against the standard approach of using C or assembly, providing valuable insights into the benefits of using Rust for this particular task.

## 2 Description of the problem

When it comes to the problem that we are trying to solve there are two main points: efficiency and automation.

### 2.1 Efficiency

In terms of efficiency at run time, natively compiled machine code will always outperform the interpreted version since the virtual machine provides an additional layer of abstraction to the actual hardware. So why would we even care about writing efficient interpreters and not just use native code compilers? There are multiple reasons why we want to care about the performance of interpreters. The main benefit we gain from using interpreters instead of compilers is that compilers are much more complex to develop and can be difficult to maintain. [4] Another reason for the interpreting approach lays in its portability. When compiling to native machine code, we limit ourself to just one target architecture. The virtual machine interpreter can execute the byte code that was generated for it on every system where the virtual machine can run and by that is not dependent on the underlying machine.

Summed up, the first problem is that interpreters will never be as fast as natively compiled machine code but by contributing to the increase of efficiency we get all the benefits of using interpreters while minimizing the disadvantages to compilers.

### 2.2 Automation

Many programmers will have the idea of implementing their own programming language at some point in their careers. Most of them will have noticed that building an interpreter is a challenging task and requires a lot of work and a clear structure. In addition to that it shows that many parts of a VM interpreter are similar and repetitive. For example the code for executing VM instructions will be similar for most of the instructions. [6] But what happens when the interpreter does not give the expected outcome in terms of efficiency? It results in manual rewriting of a codebase just to change the implementation of some part of the VM interpreter to see if the performance increases. Rewriting the whole interpreter to test if the performance is better using a different development approach is not only time-consuming but also error-prone. One solution

to this problem is automation. The developer should be able to provide a configuration file and based on that we will generate an efficient VM interpreter. It will already use efficient implementation techniques and come with built-in optimizations. It also provides easy extensibility for the user without needing to change any source code.

### 3 Literature review

This literature review compiles a selection of influential papers that explore various aspects of efficient interpreter design or the generation of them. By analyzing these papers, we aim to highlight different approaches and concepts used to develop efficient interpreters and analyze them regarding complexity. We'll cover the choice of programming language, multiple ways the interpreter can be implemented, types of virtual machines, and the output format of the compiler.

#### 3.1 Programming Language

Most of the code written, targeting the problem of efficient interpreters or compilers, such as the referenced papers [6], [3], [11] and [5] is either written in C or directly in assembly language. But why are developers using languages such as C when it is known to be error prone or even assembly which is known to be even more complex and hard to maintain. The reason for this is abstraction. Languages such as Python or Java are so easy to write because the language itself provides many layers of abstraction from what is actually happening on the machine level. This makes it easier for the programmer to write safe code. The downside of this abstraction is that it limits the control the programmer has over the system. [9] In compiler development, there are many situations where it is crucial for the programmer to be able to use little tricks on the lowest layers of the system to reach a maximum performance gain. In the paper "Optimizing an ANSI C Interpreter with Superoperators" the author Todd A. illustrates the superiority of assembly by writing:

"The interpreter is implemented in assembly language. Assembly language enables important optimizations like keeping the evaluation stack pointer and interpreter program counter in hardware registers." [11]

Another example of an optimization that is complicated to implement in high level languages such as Python but easy to implement in assembly and in some versions of C is the threaded code or computed goto optimization.

As a limitation to C the authors of 'vmgen - A Generator of Efficient Virtual Machine Interpreters' state:

"This technique cannot be implemented in ANSI C, but it can be implemented in GNU C using the labels-as-values extension." [6]



### 3.2 Interpreter approaches

There are several ways to implement an interpreter. We are looking at the 2 most widely used practices. Both of them start in the parsing phase by creating an Abstract Syntax Tree (AST). The tree consists of nodes each representing a construct in the programming language, for example, a statement or an expression. The tree walk interpreter is the simplest way to create an interpreter when the AST is already built. As the name suggests it traverses the tree and when it encounters a node it executes the corresponding operation. Tree walk interpreters are easy from the viewing point of the implementor. The tree closely mirrors the structure of the source code and by that makes errors easier to debug in the tree structure. [1] This simplicity comes with the cost of reduced performance. One reason why a tree-walk interpreter is slower than a bytecode interpreter is that as they traverse the tree they perform execution for each node which can result in redundant operations. Furthermore, it is hard to use optimizations, since it directly executes the nodes. [1]

Bytecode interpreters are traversing the AST and building up a sequence of instructions which then will be executed. The step of creating an intermediate representation (IR) opens up a new way to operate on the code. Now that we work with a series of instructions instead of a tree of nodes we can perform all sorts of optimizations on it. This makes the bytecode interpreter more efficient. It also makes it more complex since now we have to generate bytecode, implement optimizations, and execute the bytecode instructions. Debugging is also more complex since the generated instructions do not mirror the code structure anymore. [1]

### 3.3 Virtual Machines

Virtual machines create an abstraction layer between the host machine and the code that is executed. [7] This enables machine independence for code written in programming languages that do not run on a certain architecture.

The paper 'An introduction to the UCSD PASCAL system' describes the development of an early virtual machine, pseudomachine, to enable the use of high-level programming languages such as FORTRAN on unsupported systems. During that time it was common for new architectures to support popular languages like FORTRAN but even then, there were many problems when transporting programs from one system to another and smaller machines had no support for high-level languages. [2]

When talking about virtual machines we only consider the 2 main paradigms in this paper, namely stack-based and register-based virtual machines. While virtual machines are great to use in combination with interpreters each architecture has its advantages and disadvantages.

Stack-based virtual machines use a stack to manage the state. [7] The stack pointer, often abbreviated with `sp`, is used to keep track of the top of the stack. Instructions like `PUSH` and `POP` can manipulate the stack by adding or removing values from it while incrementing or decrementing the stack pointer. The main advantage of a stack-based virtual machine lies in its simplicity. The implementation is straightforward, the instructions are kept simple and optimizations are easy to implement. [7]

Register-based virtual machines work on a different principle. They use a set of virtual registers which are designed to mimic hardware registers as a place to store data. Instructions such as `LOAD` or `STORE` can modify the registers. [7] The instructions used in register-based virtual machines are larger since they need to encode the operands. As an example, the `STORE` instruction needs to know the value to store and the register where it should store the value. [7]

The execution cycle involves 3 steps. Fetching, decoding, and executing. First, the machine needs to get the instruction from memory. Then it needs to decode the instruction to know which operation to perform and which registers to use. This decoding step can add a small overhead to the machine but since it usually involves just simple logical operations it is still very fast. [3] In the last step, the operation is executed and the registers are modified accordingly. Since all operations are performed on registers, the need to push and pop from the stack is eliminated which results in efficient data access. [6]

Regarding the effort it takes to create an efficient register-based virtual machine interpreter M. Anton Ertl says:

”From the view of the compiler writer, many languages can be easily compiled for stack machine code. To achieve better performance with a register machine, the compiler must perform optimizations, e.g., global register allocation (which needs data flow analysis). This would eliminate one of the advantages of using an interpreter, namely simplicity.” [5]

A good example to compare stack-based virtual machines with register-based machines on a real-world project is Lua. Lua runs its programs on a virtual machine interpreter and therefore compiles the code into instructions for a virtual machine.

When Lua was first created in 1993 it used a stack-based virtual machine. 10 years later in 2003 with the release of Lua 5.0 they changed their approach and now use a register-based VM. [8] To understand how Lua profited from the register-based virtual machine we first need to understand how Lua works. Lua creates a prototype for each function in the code. The prototype contains an array with all the instructions to execute this function and an array of all the Lua values and constants used. [8] Upon entering a function, an activation

record is preallocated on the stack. This record holds all function registers, where local variables are stored. Resulting in efficient variable access. [8] In addition to that register-based instructions do not need any PUSH or POP instruction which are expensive in Lua because of the way Lua represents values.

Looking at the disadvantages of register-based VMs in terms of generated code size and single instruction size in Lua we can see that they are not too bad. Lua's stack-based virtual machine had an instruction set of 49 instructions where the register-based approach only needed 35. The size of a single instruction increased from 2 bytes to 4 bytes with the change of the virtual machine architecture but this relativizes itself due to less generated instructions.

Here an example:

```
local a, t, i
a = a + i
a = a + 1
a = t[i]
```

Figure 1: *Example Lua code from [8]*

The generated instructions vary a lot from Lua 4.0 to Lua 5.0.

PUSHNIL	3
GETLOCAL	0
GETLOCAL	2
ADD	
SETLOCAL	0
GETLOCAL	0
ADDI	1
SETLOCAL	0
GETLOCAL	1
GETINDEXED	2
SETLOCAL	0

Figure 2: *Generated instructions for Lua 4.0 (stack-based VM) from [8]*

Comparing Figure 2 and Figure 3 we can see even though each instruction in Figure 3 has 3 operands, resulting in a larger size, it just generated a total of 4 instructions compared to 11 in Figure 2.

LOADNIL	0 2 0
ADD	0 0 2
ADD	0 0 250 ;1
GETTABLE	0 1 2

Figure 3: *Generated instructions for Lua 5.0 (register-based VM) from [8]*

### 3.4 Output Format

The frontend of the interpreter, the compiler, generates output. What kind of output depends on the use case. For example, for virtual machine interpreters the compiler outputs byte code that can be run by the virtual machine. This is platform-independent but comes with the cost of reduced performance. A native-code compiler outputs machine code for a specific architecture. This is faster than interpreting but is limited when it comes to portability. [6]

Another approach is to compile to C. With that, we have portability and gain all the optimization that were added to the C compiler in the last 30 years. Some disadvantages are stated by the authors of 'The Structure and Performance of Efficient Interpreters':

"Why not write a compiler to C? While this method addresses the ease-of-implementation and retargeting issues, it is not acceptable when interactivity, quick turnarounds, or non-C features (e.g., run-time code-generation or guaranteed tail-call optimization) are required." [4]

The paper 'Optimizing an ANSI C Interpreter with Superoperators' shows a hybrid approach that generates "efficient code that includes a small amount of native code with interpreted code". [11]

"This hybrid approach allows the object files to maintain all C calling conventions so that they may be freely mixed with natively compiled object files. The interpreted object code is approximately the same size as equivalent native code, and runs only 3-16 times slower." [11]

## 4 Approaches

In the following section, we will go into detail about why we have used the techniques and tools as we did, re-evaluate the decisions and present learnings. The choice of the implementation language had probably the most influence on the efficiency of our virtual machine interpreter. The industry standard for implementing such systems is either C or assembly. We made the decision to use Rust early on to explore how we can use its higher-level language features,

such as its enums, to do low-level programming and see if we can get close to the efficiency of C implementations with the convenience that Rust gives us during development. Rust prevents many errors from happening during development by for example enforcing strict rules on ownership of data and providing high-level language features such as iterators, match statements and a rich standard library of algorithms and containers such as HashMaps. To represent byte code in Rust we decided to use its powerful enums. This is a very convenient way since enums can have fields attached to certain variants and are an actual type not like in C where enums are just numbers. While this approach is very convenient it is not very efficient. One byte code instruction as we designed the enum takes up 32 bits of memory. The largest enum variant is the jump instruction as it holds a string called 'label' and an i32 called 'offset'. Since Rust's enums behave like a tagged union every instruction has the size of the largest instruction.

A string in Rust contains:

- 8 bits -> pointer to the chars
- 8 bits -> len of the string
- 4 bits -> the i32 value
- 8 bits -> enum tag
- + padding

adds up to 32 bits.

A more efficient approach to represent byte code in Rust would be to just define each instruction by a unique integer. By doing this, we would basically use the same technique as using enums in C and lose the convenience of using Rust's language features. A big downside that followed the decision to use Rust was the way we implemented the computed goto optimization. The detailed implementation is discussed in section 4. Since Rust does not allow the usage of labels like C does, we had to use an alternative technique where we stored the function pointers to the routines that execute a byte code instruction, in a hash map and had to use the discriminants as an index. This additional indirection at the end of every executed instruction caused an extreme performance slowdown. The typical approach in C would be to store the addresses of the labels to the routines in an array and index it by using the C enum representation of an instruction. Earlier we evaluated that register-based virtual machines, when set up correctly, can be faster and more efficient than stack-based VMs. In our implementation, we still decided to use a stack-based virtual machine. The reason for that lies in its simplicity. This project aims to explore how to build a virtual machine interpreter in Rust. Therefore we tried to keep the implementation as simple as possible. The possible implementation of a register-based virtual machine in the future is described in Section 8.

## 5 Description of work

The goal of this project is to build a virtual machine interpreter and explore the implementation of such in the Rust programming language.

### 5.1 Tools used

During the development of this project, we used three programming languages.

1. **Rust:**  
We used Rust to build all parts of the virtual machine interpreter.
2. **Python:**  
We used Python to capture and visualize benchmark results using matplotlib.
3. **Imp:**  
Imp is the input language the VM interpreter operates on. 5.5

### 5.2 Virtual Machine

The virtual machine, the backend of the VM interpreter, is the part that executes the byte code that was generated by the compiler.

#### 5.2.1 Implementation

The virtual machine we have built is very basic. It contains a representation of the stack, the program counter (pc) which is used as an index into the series of instructions and a HashMap to hold the global variables.

```
pub struct ByteCodeInterpreter {  
    stack: Vec<usize>,  
    pc: i32,  
    variables: HashMap<String, usize>,  
}
```

*ByteCodeInterpreter struct definition*

The optimized version of it implementing the ‘computed goto’ optimization (see section 4) is a bit more complicated. Here we also have the mapping of ByteCode to the function that executes this bytecode. This optimization is described in more detail at 4.

```
pub struct ByteCodeInterpreterThreaded {  
    stack: Vec<usize>,  
    pc: i32,  
    variables: HashMap<String, usize>,  
    ops: HashMap<Discriminant<ByteCode>, Instruction>,  
}
```

```

        instructions: Vec<ByteCode>,
    }

```

*ByteCodeInterpreterThreaded struct definition*

### 5.2.2 Stack

The virtual machine we built is stack-based. This means there are no registers and all the data is handled on the stack. When you add two numbers there will be a PUSH instruction for both of the values and an ADD instruction which will pop the last two values from the stack, adds them together and pushes the result back on the stack.

The code `'1 + 2'` would produce the following instructions:

```

inst:    stack:
         []
push 1   [1]
push 2   [1, 2]
add      [3]

```

### 5.2.3 Byte code

There are many ways to model a byte code in Rust. We have decided to use Rust's powerful enums. They are not only able to perfectly represent each instruction with its parameters but also helps during development due to Rust's powerful type system.

```

#[derive(Debug, Clone, PartialEq)]
pub enum ByteCode {
    Push(usize),
    Pop,
    Add,
    Sub,
    Mul,
    Var(String),
    Eq,
    NEq,
    Lt,
    Gt,
    Lte,
    Jz {
        label: String,
        offset: i32,
    },
}

```

```

    ...
}

```

*Rust representation of the Byte Code*

## 5.3 Interpreter

The interpreter consists of a scanner and a parser. We have decided to handwrite both parts instead of using tools like 'Lex' which can be used to automatically generate a scanner from a definition file. Even though the scanner and the parser are not the most performance-critical parts of an VM interpreter we decided to handwrite them in Rust to have full control over what is happening.

### 5.3.1 Scanner

The scanner was implemented very straight forward by using a match statement inside a while loop to iterate the source code and output a stream of tokens. Tokens are defined as the token type and an optional literal.

```

pub struct Token {
    pub token_type: TokenType,
    pub literal: Option<Object>,
}

```

*Token definition*

The 'Object' type represents a literal expression such as a number, an identifier for a variable or a boolean value like True or False. Other than that it is used to carry information about errors when they occur.

```

pub enum Object {
    Num(f64),
    Bool(bool),
    Variable(String),
    DivByZeroError,
    ArithmeticError,
}

```

*Object definition*

```

pub enum TokenType {
    LeftParen,
    RightParen,

    Minus,
    Plus,
}

```



```

    ...
    Eof,
}

```

*TokenType definition*

### 5.3.2 Parser

The parser takes the stream of tokens as input and builds up an Abstract Syntax Tree (AST). In the Rust code, it is represented as a `Vec<Rc<Stmt>>`.

As an example, the code `'x := 10;'` will get passed into the scanner. It returns a token stream like this: `'[Identifier, Assignment, NumberLiteral, Semicolon, EOF]'`. The parser will iterate through the list of tokens and builds up a tree structure (Abstract Syntax Tree) like this:

```

ExpressionStmt {
    expr: AssignExpr {
        name: "x",
        value: LiteralExpr {
            value: Num(10)
        }
    }
}

```

### 5.3.3 Interpreter (testing)

To be able to test the scanner and parser before writing the bytecode generator we decided to build a small tree walk interpreter that just runs the code. It uses the visitor pattern to traverse the abstract syntax tree generated by the parser and executes the code immediately.

To use the visitor pattern approach we have implemented a statement visitor and an expression visitor.

```

pub trait StmtVisitor<T> {
    fn visit_block_stmt(&self, stmt: &BlockStmt) -> Result<T, ()>;
    fn visit_if_stmt(&self, stmt: &IfStmt) -> Result<T, ()>;
    fn visit_expression_stmt(&self, stmt: &ExpressionStmt) -> Result<T, ()>;
    fn visit_print_stmt(&self, stmt: &PrintStmt) -> Result<T, ()>;
    fn visit_while_stmt(&self, stmt: &WhileStmt) -> Result<T, ()>;
}

```

*Statement visitor example*

For each statement, we have set up a structure holding the necessary data. For example in the structure for the if statement, we need the condition, which is of type `Expr`, the statement in the then\_branch and the optional statement for the else branch.

```
#[derive(Debug)]
pub struct IfStmt {
    pub condition: Rc<Expr>,
    pub then_branch: Rc<Stmt>,
    pub else_branch: Option<Rc<Stmt>>,
}
```

*Structure holding data for an if statement*

## 5.4 Optimizations

Optimizations play a critical role in the field of compiler design. From the endless list of optimizations we chose a couple interesting ones, that are not too hard to implement, to dive into detail about and use in our virtual machine interpreter.

Some well known optimization techniques are:

### 1. Dead Code Elimination:

Remove parts of the code that are never to be used.

Consider the following C code. In this case, the variable ‘c’ is never used. So the statement ‘int c = 3;’ does not need to be compiled.

```
int main() {
    int a = 1;
    int b = 2;
    int c = 3; // never used

    return a + b;
}
```

### 2. Constant folding:

When calculations are only using constant values and by that can be evaluated during compile time the compiler will directly replace it with the computed value.

Consider the following C code. In this example ‘int a = 10 + 20 + 30;’ only uses values known at compile time. So instead of multiple PUSH instructions and ADD instructions the compiler could internally replace the statement with ‘int a = 60;’ and by that just have one value.

```
int main() {
    int a = 10 + 20 + 30;
    printf("%d", a);

    return 0;
}
```

### 3. Superinstructions:

Superinstructions is a term used for the case when we can combine two or more instructions into a single big instruction.

Consider the following C code.

```
int main() {
    int i = 0;
    while (i < 10) {
        i += 1;
    }
}
```

If we naively 'compile' this code we would get instructions like this:

```
// in the loop
load i -- get the variable at i and push it to the stack
push 1 -- push the value '1' to the stack
add     -- pop the first 2 values from the stack,
          add them, and push the result back on the stack
```

We always need 2 instructions to push the constant value to the stack and then add the top 2 values on the stack together.

But we can create a superinstruction called 'push\_add' which would change the generated code to this.

```
// in the loop
load i      -- get the variable at i and push it to the stack
push_add 1 -- superinstruction that does the push and then the add
```

That change might seem very insignificant but we can extend this idea of superinstructions to combine already built superinstructions with each other and by that save many instructions.

### 4. Computed goto:

When executing the byte code inside of a virtual machine interpreter, for example, you will usually find a big switch statement inspecting the current instruction and executing it. In our case, since we are using Rust we have a big match statement like this:

```

while self.pc < instructions.len() as i32 {
    let inst = &instructions[self.pc as usize];
    match inst {
        ByteCode::Push(value) => {
            self.stack.push(*value);
        }
        ByteCode::Pop => {
            self.stack.pop().unwrap();
        }
        ByteCode::Add => {
            let a = self.stack.pop().unwrap();
            let b = self.stack.pop().unwrap();
            self.stack.push(b + a);
        }
        ...
    }
}

```

This ‘while’ loop will iterate over every single instruction in the program. That means for every instruction it has to go through all the possible values inside of the ‘match’ statement until it finds the matching result.

The ‘Computed goto’ optimization tries to make this process more efficient. To implement that, we have to create an array of function pointers where the index of the array corresponds to an instruction or use a HashMap with the byte code instruction as the key and the function as a value like we did in Rust.

```

pub type Instruction = fn(interp: &mut ByteCodeInterpreterThreaded);
... {
    ops: HashMap<Discriminant<ByteCode>, Instruction>
}

ops.insert(std::mem::discriminant(&ByteCode::Push(0)), Self::op_push);
ops.insert(std::mem::discriminant(&ByteCode::Pop), Self::op_pop);
ops.insert(std::mem::discriminant(&ByteCode::Add), Self::op_add);

```

Then each of the functions (Self::op\_push, Self::op\_add, ...) will finish with a call to a ‘next()’ function. This next() function will increment the program counter (the current position in the byte code) and call the function which is mapped to the next instruction in the HashMap. In our case in Rust, the next function looks like the following.

```

fn next(&mut self) {
    // incrementing the program counter;
    self.pc += 1

    // check for end of stream
    if self.pc >= self.instructions.len() as i32 {
        return;
    }

    // call the next function
    self.ops[&self.instructions[self.pc as usize]](self);
}

```

This threaded code approach results in a short and fast instruction dispatch sequence and can lead to better branch prediction accuracy on machines with branch target buffers as stated by the authors of vmgen [6].

#### 5.4.1 Replacing superinstructions in the byte code

To generate a byte code that uses superinstructions we first need to generate the normal byte code. In a second iteration we need to go through the sequence of instructions and check if we can combine any of the instructions into a superinstruction. During this process we need to iterate through all the instructions once and by that the time it takes grows linear with the size of the byte code. This approach of applying the optimization is also called peephole optimization as we are scanning through the list of instructions only looking at a small part of it at a time.

### 5.5 Input language

When you want to build a virtual machine interpreter you need to have an input language. We have decided to use an already existing programming language instead of creating our own since this project aims to explore the implementation of a VM interpreter and not the design of a programming language. We chose to use the programming language Imp. [10]

"Imp is a simple imperative programming language which embodies a tiny core fragment of conventional mainstream languages such as C or Java". [10]

```

Z := X;
Y := 1;
while Z != 0 do
    Y := Y * Z;
    Z := Z - 1;
end

```

### *Example code in the Imp programming language*

Not only is Imp simple but it is also already fully defined and has many examples we can test against. The book ‘IMP Simple Imperative Programs’ by Benjamin C. Pierce [10] describes the language in detail and provides additional information such as BNF grammar definitions of the syntax. This fact makes it easy to understand and provides a guideline when developing the scanner and parser of the interpreter.

The language has no local scope resulting in all variables falling into global scope. This fact makes it easier for us since we do not need to keep track of scopes or environments and the variable lookup in our interpreter can be implemented using just a single HashMap that holds all global variables.

To declare variables in Imp we do not need to provide an explicit type. Since all variables in Imp are numbers it makes it very convenient for us to implement.

```
a := 1;
b := a;
```

### *Variable assignment in the Imp programming language*

Imp defines while loops which are delimited by the ‘do’ and ‘end’ keywords and does not need any parentheses around the condition. Those particular circumstances make Imp really easy to parse.

```
a := 0;
while a < 10 do
  a := a + 1;
end
```

### *While statement in the Imp programming language*

Similar to while loops, if statements are also defined without any unnecessary parentheses and are delimited by describing keywords like ‘then’, ‘else’ and ‘end’.

```
a := 0;
if a < 10 then
  a := a + 1;
end
```

### *If statement in the Imp programming language*

For convenience and testing purposes we added our own keyword ‘print’ to the language. It helps us to get an inside view of the execution of the interpreter, which can be very helpful during development or while debugging.

```

a := 0;
if a < 10 then
    print a + 1;
else
    print a - 1;
end

```

*Print statement in the Imp programming language*

Furthermore, we have implemented the modulo operator and a ‘continue’ statement, to manually continue to the next iteration of a loop. By doing that, we have more possibilities to create benchmark programs. Imp does not include function calls or any of the commonly known features of today’s high-level languages such as classes, lambdas, structs or even strings and arrays.

## 6 Results & Benchmarks

### 6.1 Setup

To create the benchmarks we have used a system with the following specs.

1. **CPU:**  
11th Gen Intel i7-11370H (8) @ 4.800GHz
2. **Operating System:**  
Fedora Linux 37 (Workstation Edition) x86\_64
3. **Rust Compiler:**  
cargo 1.66.0 rustc 1.66.0

The script ‘benchmark.py’ runs the virtual machine interpreter using every program defined in the ‘./benchmarks’ directory as an input. It then generates visualization using the python library matplotlib. To collect the benchmark timings and reduce the chance of outliers, that can occur because of external reasons such as system load or memory availability, the script executes each command 10 times and calculates the average. To be able to specify which superinstructions should be used in the current execution of the VM interpreter we have used Cargo’s ‘features’ option.

The following code shows how to write conditional code that will be turned on or off during compilation in Rust.

```

enum ByteCode {
    ...
    #[cfg(feature = "PushAdd")]
    PushAdd(usize),
    #[cfg(feature = "AssignPushAdd")]

```

```

        AssignPushAdd {
            name: String,
            value: usize,
        },
    },
}

```

## 6.2 Test Programs

To get the benchmark data we used multiple small programs that generate different byte code. By doing that, we can see which optimization gives the most benefit for what kind of program.

As test programs, we used:

### 1. **Factorial.imp:**

A simple program to calculate 10 factorial.

```

result := 1;
n := 10;
while n > 1 do
    result := result * n;
    n := n - 1;
end
print result;

```

### 2. **Fib\_70.imp:**

A program to calculate and print the first 70 fibonacci numbers.

```

a := 0;
b := 1;
n := 70;

while n > 1 do
    print a;
    temp := a;
    a := b;
    b := temp + b;
    n := n - 1;
end

```

### 3. **Gcd.imp:**

A program to calculate the greatest common divisor (GCD) of two numbers.

```

a := 234234;
b := 234;

while a != b do

```



```

        if a > b then
            a := a - b;
        else
            b := b - a;
        end
    end

    print a;

```

4. **If.imp:**

A program that uses many if statements and assignments inside of a loop.

```

i := 1;
a := 0;

while i < 1000 do
    i := i + 1;
    a := a + 1;

    if a < i then
        a := 0;
    end
    if a < i then
        a := 0;
    end
    if a < i then
        a := 0;
    end
    if a < i then
        a := 0;
    end
end
end

```

5. **Increment\_loop.imp:**

A program that uses many assignments and additions.

```

a := 0;
b := 0;
c := 0;
d := 0;
e := 0;
f := 0;
g := 0;
i := 0;

while i < 1000 do

```

```

        i := i + 1;
        a := a + 1;
        b := b + 1;
        c := c + 1;
        d := d + 1;
        e := e + 1;
        f := f + 1;
        g := g + 1;
    end

    print i;

```

6. **Sum.imp:**

A program that calculates the sum of the first 1000 natural numbers.

```

N := 1000;
Sum := 0;
i := 1;

while i <= N do
    Sum := Sum + i;
    i := i + 1;
end

print Sum;

```

7. **Fizz\_buzz.imp:**

A program that uses the modulo operator to calculate the famous fizzbuzz problem.

```

i := 0;

while i < 1000 do
    i := i + 1;
    if i % 3 == 0 && i % 5 == 0 then
        print 0;
        continue;
    end
    if i % 3 == 0 then
        print 1;
        continue;
    end
    if i % 5 == 0 then
        print 2;
        continue;
    end
end

```

```

    print i;
end

```

Each of the programs will generate a byte code that uses some instructions intensively. By that, we can see which superinstruction has the most influence on run time and also see the time the generation for the byte code using this superinstruction combination took.

### 6.3 Results

The following images show the performance analysis charts.

The y-axis is calibrated in milliseconds and shows the time.

On the x-axis we show the different files. Each entry has 3 values attached to it.

1. **Generation time (blue):**  
The time it took to generate the byte code for the application.
2. **Interpreting time (orange):**  
The time it took to interpret the byte code.
3. **Interpreting time (threaded) (green):**  
The time it took to interpret the byte code with the computed goto (or threaded code) optimization enabled.




Figure 4: Benchmark results without any optimizations enabled

The first diagram 4 has no optimizations enabled. We can use this as the base line. Now we can compare figure 4 with the figures that have superinstructions enabled.




Figure 5: Benchmark results with the superinstruction PushAdd enabled




Figure 6: Benchmark results with the superinstructions PushAdd and AssignPushAdd enabled

When we compare the graphs, we can see that the threaded code optimization causes a slowdown of approximately 300% compared to the non-optimized version. But why is this optimization making the program slower? We have already explained how we have implemented the computed goto optimization for that see Section 4. There you can see that we used a hash map to store the mapping of byte code instruction to the function that implements it. The key is a `usize`. This `usize` is the result of a call to the `std::mem::discriminant` function. It returns the index of a given variant in an enum. Since Rust does not by default allow an enum with fields that hold data to be the key of a hash map we had to use this approach. After every instruction, we need to call the discriminant function to get the index of the next byte code instruction. This overhead is most likely the main reason for the slowdown.




Figure 7: Benchmark results with the superinstructions PushAdd, AssignPushAdd and PushAssign enabled

When we are comparing Figure 5 (PushAdd) and Figure 4 (No optimizations) we can see that the implementation of the superinstruction PushAdd which combines the PUSH and the ADD instruction resulted in slight improvements in the run time. The biggest performance benefit was gained in the increment\_loop.imp program 6. The generated byte code for this program uses the combination of pushing the value 1 to the stack and adding it to a variable in a loop. Combining those two instructions we can gain almost a 20% speed improvement in the threaded code and a 6% speed improvement in non-threaded code.

The following shows the generated byte code for the increment\_loop.imp program without using any superinstructions.

inside the loop...

```

Var("a")
Push(1)
Add
Assign("a")
Var("b")
Push(1)
Add
Assign("b")
...

```

This shows the generated byte code using the PushAdd instruction.

inside the loop...

```

Var("a")
PushAdd(1)
Assign("a")
Var("b")
PushAdd(1)
Assign("b")
...

```

As you can see we saved one instruction per incremented variable per iteration in the loop. To get an even bigger performance increase on the `increment_loop.imp` program we can enable two superinstructions. As we can see in Figure 6 (PushAdd, AssignPushAdd) the run time was decreased to only 32.7 ms in threaded code. This is an improvement in speed of nearly 30% in threaded code and 10% in non-threaded code. When we combine the previously built PushAdd superinstruction with the ASSIGN instruction we can reduce the number of instructions even more to just two instructions per incremented variable.

inside the loop...

```

Var("a")
AssignPushAdd { name: "a", value: 1 }
Var("b")
AssignPushAdd { name: "b", value: 1 }
...

```

Since none of the test programs results in a generated bytecode that exceeds 100 instructions the differences in generating are minimal but in some cases still notable. When we compare the generation time of `fizz_buzz.imp` in Figure 4 and Figure 7 we can see that the time it took to generate the byte code increased by 0.2ms. This increase might seem very little and irrelevant but the generated byte code only consists of 57 instructions. If this would be a big program resulting in thousands of instructions the time of code generation would noticeably increase as the time needed to create superinstructions grows linear to the byte code (see Section 5.4.1).

## 7 Summary

The efficiency of a virtual machine interpreter is not just dependent on the optimizations it implements on the generated code but also requires the virtual machine and the interpreter to use efficient techniques. Factors like the design of the byte code for the VM or layout of the byte code in the memory are important factors to consider when trying to write an efficient VM interpreter.

We explored the implementation in Rust, using its high-level features. We

showed that those features, such as Rust’s enums, are a convenient way to implement a VM interpreter but this might not be the most optimal approach to it. When comparing it to the way you would implement, for example, the byte code in C, we can see that the technique that is not that convenient for the developer is actually way faster and more efficient. The same applies to the technique used to implement optimizations like the threaded code in the interpreter.

When talking about optimizations, we showed that even simple optimizations like superinstructions can have a positive impact on run time performance while causing a slowdown during the generation of the byte code.

## 8 Future Work

The last section of this paper will focus on possible ways to improve this project in the future.

Looking at the title of this paper and comparing it to the write-up we can see that we are missing an essential part. The part of automation. The idea behind the automation part is to give the user of this program the opportunity to provide a configuration file which can define all sorts of information that then will be used for the generation of the virtual machine interpreter. Some of the information can be custom superinstructions for example. The idea is to let the user define their own superinstructions and write them down in a syntax we can understand to automatically implement them during generation.

Another way to improve this project in the future is to think about a different way to model the byte code inside of the Rust code as our chosen approach was not the most efficient.

To further improve the efficiency we can implement more optimization techniques in the future. To name a few optimizations that should be pretty easy to implement at this point are for example ‘dead code elimination’ as described in Section 1.

Another thing that can be done in the future is to add more superinstructions. For that, it would be useful to create more benchmark examples and see what instructions are used frequently to create superinstructions out of them.

An interesting extension to the integration of superinstructions was mentioned in the paper ‘Optimizing an ANSI C Interpreter with Superoperators’ [11] where they have used a heuristic method to infer good sets of superinstructions. This approach automates the process of finding good superinstructions to use.

Yet another way to further develop this project is to change the architecture of the virtual machine and use a register-based machine. Doing this would also

result in a change of the byte code representation and optimally make the VM interpreter more efficient.

Finally after making the virtual machine interpreter efficient to work with Imp, we could project this system onto a different input language. In the future we could support rather simple scripting languages like Lua or Ruby or more complex languages such as Python. Another idea would be to let the user define the syntax of the input language in a configuration file using and we generate the scanner and the parser accordingly.



## References

- [1] Nikola Bebic and Christian Humer. Operation dsl making bytecode interpreters a walk in the park. 2022.
- [2] Kenneth Bowles and James Hollan. An introduction to the ucsd pascal system. *Behavior Research Methods*, 10:531–534, 07 1978.
- [3] M. Ertl and David Gregg. The behavior of efficient virtual machine interpreters on modern architectures. 11 2001.
- [4] M. Ertl and David Gregg. The structure and performance of efficient interpreters. *J. Instruction-Level Parallelism*, 5, 11 2003.
- [5] M. Anton Ertl. Stack caching for interpreters. page 315–327, 1995.
- [6] M. Anton Ertl, David Gregg, Andreas Krall, and Bernd Paysan. Vmgen: A generator of efficient virtual machine interpreters. *Softw. Pract. Exper.*, 32(3):265–294, mar 2002.
- [7] Ruijie Fang and Siqi Liu. A performance survey on stack-based and register-based virtual machines. 2016.
- [8] Roberto Ierusalimschy, Luiz Figueiredo, and Waldemar Celes. The implementation of lua 5.0. *J. UCS*, 11:1159–1176, 01 2005.
- [9] Brijender Kahanwal. Abstraction level taxonomy of programming language frameworks. 11 2013.
- [10] Benjamin C. Pierce, Arthur Azevedo de Amorim, Chris Casinghino, Marco Gaboardi, Michael Greenberg, Cătălin Hrițcu, Vilhelm Sjöberg, and Brent Yorgey. "Logical Foundations", volume "1" of "Software Foundations". "Electronic textbook", "2023".
- [11] Todd A. Proebsting. Optimizing an ansi c interpreter with superoperators. page 322–332, 1995.